

# Bayesian Reparametrisation of Neural Networks

by Moritz Grünbauer



# Structure

- 1) Motivation
- 2) What is Uncertainty?
- 3) Bayesian Neural Network
- 4) Original Work by Bruss, et al.
- 5) Reproduction
- 6) Results
- 7) Conclusion



# Motivation

- Neural Networks (NN) are inherently uncertain in their results
- But they pretend to “have an answer for everything”
- If we’re able to quantify uncertainty, we can assign value to predictions of Nns
- Bayesian NNs are a way to do that, but can we make it easier?

# What is Uncertainty?

## Epistemic:

- Can be “trained away”
- Better Technologies and Methods can reduce this

## Aleatoric:

- Inherent to the data
- Example:
- Classifying 1 as 7 or vice versa
    - Because they share features

# Bayesian Neural Networks

- Probability distributions rather than point estimates
- Accepting only predictions with low uncertainty

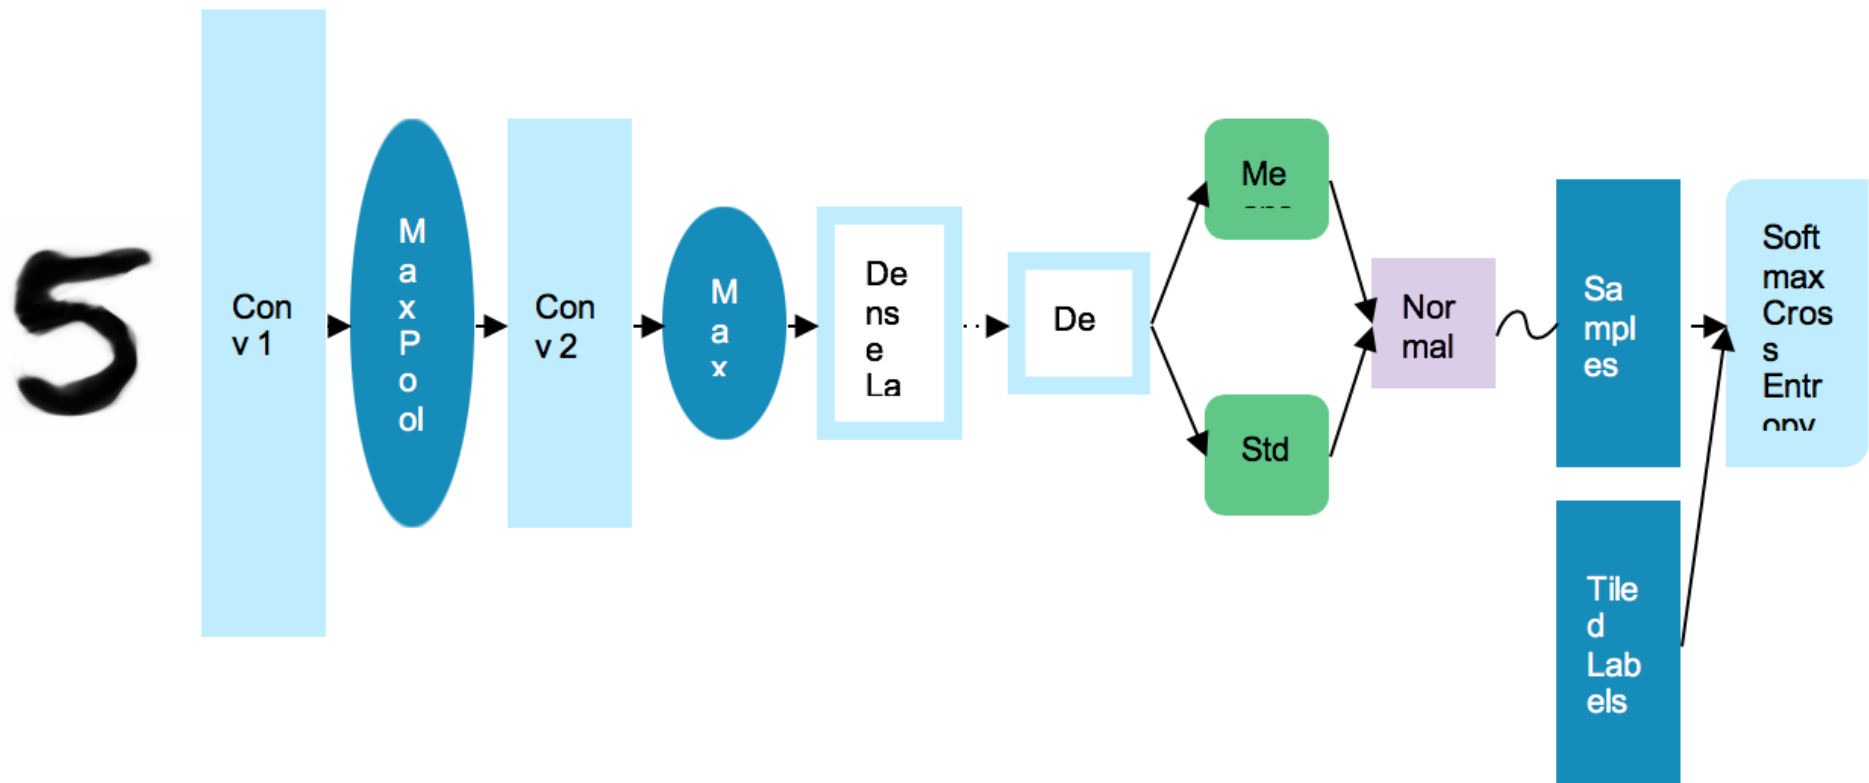
But:

- Have to carry distributions through all layers

**What if:  
We could make regular NNs  
“Bayesian”?**

# Original work by Bruss, et al.

- Added reparametrisation to regular MNIST classifier to get awareness of uncertainty



# Original work by Bruss, et al.

- They report: jump from 97% to 99.3% accuracy
- No mention of accepting only results with a minimum certainty



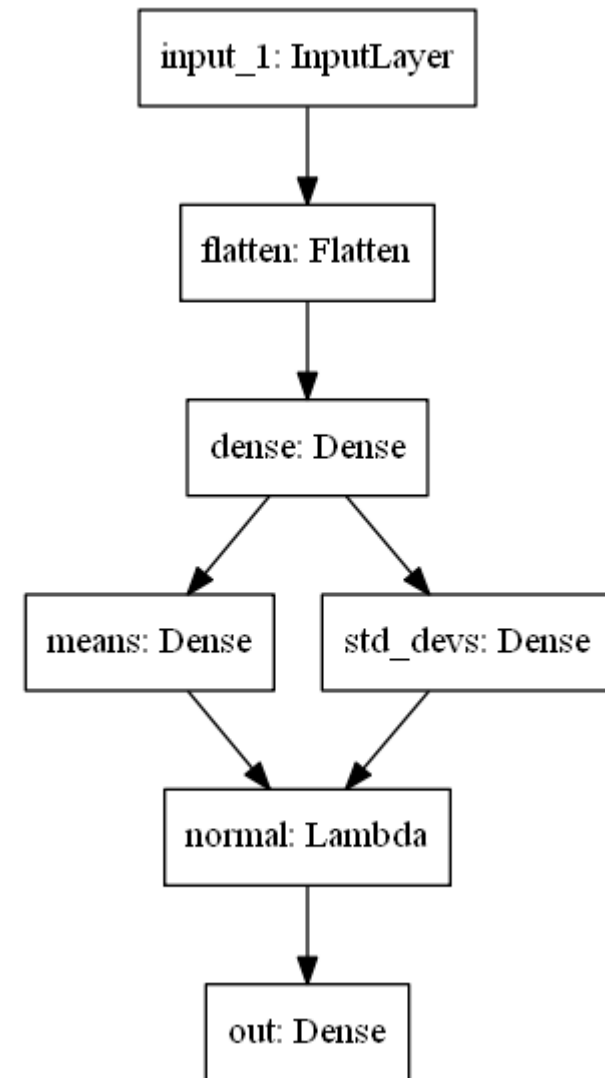


# Reproduction

- No full source code given, except:

```
80 ##### Fit a Gaussian over dense layer output #####
81
82 # Means of Gaussian (10 classes)
83 locs = tf.layers.dense(inputs=dense2, units=10, name="means")
84
85 # Standard Deviations of Gaussian (10 classes)
86 scales = tf.layers.dense(inputs=dense2, units=10,
87 | | | | | | | | | | name="std_devs", activation=tf.nn.softplus)
88
89 # Parameterize the Gaussian
90 dist = Normal(loc=locs, scale=scales)
91
92 # Sample from Gaussian 1000 times
93 num_sample = 1000
94 logits = dist.sample([num_sample], name='logits')
95
96 # Change shape of sampled logits
97 logits = tf.transpose(logits, [1, 0, 2])
98
99 # Replicate the true label 1000 times, once for each sample
100 labels = tf.tile(labels[:, tf.newaxis], [1, num_sample])
```

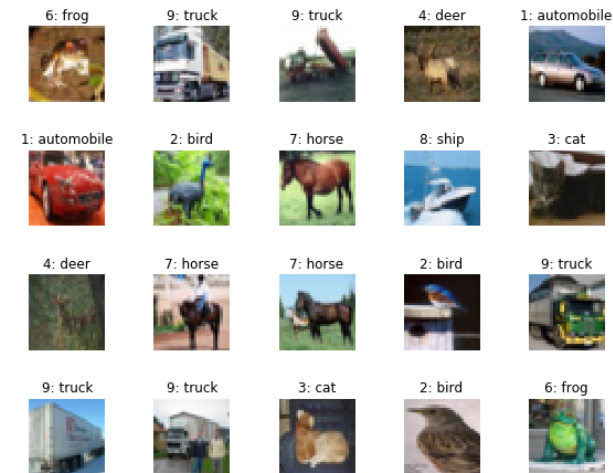
# Reproduction

- Final version of reproduction uses:
  - Tensorflow
  - Keras
  - Lambda Layer
- Basis is an official Tensorflow MNIST tutorial



# Results

- On MNIST dataset:
  - No increase in accuracy
  - Significant increase in training / testing time (4x longer)
  - The 1000 samples are basically just 1000 repeats of point estimate with tiny deviation
- On CIFAR-10 dataset:
  - Same results as MNIST



# Conclusion

- Accuracy increase was not reproducible
- Maybe reproduction was wrong due to sparse information

**It was not possible to gain the advantages of Bayesian NNs with reparametrisation.**